

```
df=pd.read_csv('/content/drive/MyDrive/insurance.csv')
df.head()
```

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	co
0	6762810145	42491	5	2.50	3650	9050	2.0	0	4	
1	6762810635	42491	4	2.50	2920	4000	1.5	0	0	
2	6762810998	42491	5	2.75	2910	9480	1.5	0	0	
3	6762812605	42491	4	2.50	3310	42998	2.0	0	0	
4	6762812919	42491	3	2.00	2710	4500	1.5	0	0	

5 rows × 23 columns

```
#import all the libraries
import pandas as pd
import numpy as np
import tensorflow as tf

df.shape

(14620, 23)

#checking null values
df.isnull().sum()
```

id	0
Date	0
number of bedrooms	0
number of bathrooms	0
living area	0
lot area	0
number of floors	0
waterfront present	0
number of views	0
condition of the house	0
grade of the house	0
Area of the house(excluding basement)	0
Area of the basement	0
Built Year	0
Renovation Year	0
Postal Code	0
Latitude	0
Longitude	0
living_area_renov	0
lot_area_renov	0
Number of schools nearby	0
Distance from the airport	0
Price	0
dtype: int64	

```
#checking the datatypes
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14620 entries, 0 to 14619
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    14620 non-null  int64
1   Date                                14620 non-null  int64
2   number of bedrooms                  14620 non-null  int64
3   number of bathrooms                 14620 non-null  float64
4   living area                         14620 non-null  int64
5   lot area                           14620 non-null  int64
6   number of floors                    14620 non-null  float64
7   waterfront present                  14620 non-null  int64
8   number of views                     14620 non-null  int64
9   condition of the house              14620 non-null  int64
10  grade of the house                  14620 non-null  int64
11  Area of the house(excluding basement) 14620 non-null  int64
```

```

12 Area of the basement      14620 non-null int64
13 Built Year                14620 non-null int64
14 Renovation Year           14620 non-null int64
15 Postal Code               14620 non-null int64
16 Lattitude                 14620 non-null float64
17 Longitude                 14620 non-null float64
18 living_area_renov         14620 non-null int64
19 lot_area_renov            14620 non-null int64
20 Number of schools nearby   14620 non-null int64
21 Distance from the airport 14620 non-null int64
22 Price                     14620 non-null int64
dtypes: float64(4), int64(19)
memory usage: 2.6 MB

```

```

from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()

```

```

#convert the string type to int type
df['Price']=le.fit_transform(df['Price'])

```

```

#split the data independent& dependent)
x=df.iloc[:,0:-1].values
x

```

```

array([[6.76281014e+09, 4.24910000e+04, 5.00000000e+00, ...,
        5.40000000e+03, 2.00000000e+00, 5.80000000e+01],
       [6.76281064e+09, 4.24910000e+04, 4.00000000e+00, ...,
        4.00000000e+03, 2.00000000e+00, 5.10000000e+01],
       [6.76281100e+09, 4.24910000e+04, 5.00000000e+00, ...,
        6.60000000e+03, 1.00000000e+00, 5.30000000e+01],
       ...,
       [6.76283062e+09, 4.27340000e+04, 2.00000000e+00, ...,
        6.12000000e+03, 2.00000000e+00, 6.40000000e+01],
       [6.76283071e+09, 4.27340000e+04, 4.00000000e+00, ...,
        6.63100000e+03, 3.00000000e+00, 5.40000000e+01],
       [6.76283146e+09, 4.27340000e+04, 3.00000000e+00, ...,
        3.48000000e+03, 2.00000000e+00, 5.50000000e+01]])

```

```

y=df.iloc[:,4:5].values

```

```

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)

```

```

from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)
print(x_train)

```

```

[[-1.46639097 -1.16155586  0.65582076 ... -0.17007044 -1.23664001
  -0.99104498]
 [-0.32364482 -1.08747669 -0.39868565 ... -0.25454935  1.2121438
  -0.99104498]
 [ 1.57399277  0.67560763 -0.39868565 ... -0.13110387  1.2121438
   0.91013359]
 ...
 [ 0.71357026  0.55708096 -0.39868565 ... -0.19790921  1.2121438
   0.12729536]
 [ 0.61620624  0.79413431  1.71032717 ... -0.18693542  1.2121438
  -0.43187481]
 [-0.46744892 -1.01339751  0.65582076 ... -0.27934626 -0.01224811
   0.46279746]]

```

```

# one-encode the geograpy column
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), [1])], remainder='passthrough')
x = np.array(ct.fit_transform(x))
print(x)

```

```

(0, 0)      1.0
(0, 241)    6762810145.0
(0, 242)    5.0
(0, 243)    2.5

```

```

(0, 244)    3650.0
(0, 245)    9050.0
(0, 246)     2.0
(0, 248)     4.0
(0, 249)     5.0
(0, 250)    10.0
(0, 251)   3370.0
(0, 252)    280.0
(0, 253)   1921.0
(0, 255)  122003.0
(0, 256)   52.8645
(0, 257)  -114.557
(0, 258)   2880.0
(0, 259)   5400.0
(0, 260)     2.0
(0, 261)    58.0
(1, 0)       1.0
(1, 241)  6762810635.0
(1, 242)     4.0
(1, 243)     2.5
(1, 244)   2920.0
:           :
(14618, 256) 52.7157
(14618, 257) -114.411
(14618, 258) 1420.0
(14618, 259) 6631.0
(14618, 260) 3.0
(14618, 261) 54.0
(14619, 240) 1.0
(14619, 241) 6762831463.0
(14619, 242) 3.0
(14619, 243) 1.0
(14619, 244) 900.0
(14619, 245) 4770.0
(14619, 246) 1.0
(14619, 249) 3.0
(14619, 250) 6.0
(14619, 251) 900.0
(14619, 253) 1969.0
(14619, 254) 2009.0
(14619, 255) 122018.0
(14619, 256) 52.5338
(14619, 257) -114.552
(14619, 258) 900.0
(14619, 259) 3480.0
(14619, 260) 2.0
(14619, 261) 55.0

```

## ▼ Build the ANN Model

```

import keras
from keras.models import Sequential
from keras.layers import Dense

```

```

#initializing the Ann
ann=tf.keras.models.Sequential()

```

## ▼ Input layer

```

# add input layer
ann.add(tf.keras.layers.Dense(units=12,activation='relu',input_shape=x_train[0].shape))

```

## ▼ Hidden Layer

```

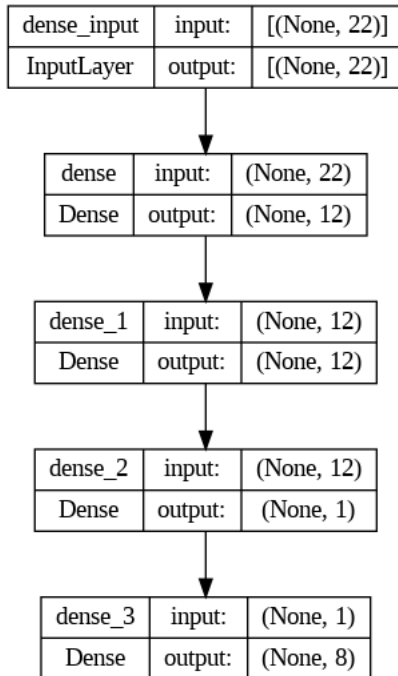
#hidding two layers
ann.add(tf.keras.layers.Dense(units=8, activation='relu'))

```

## ▼ Output Layer

```
ann.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))
```

```
from tensorflow.keras.utils import plot_model
plot_model(ann,
            to_file="model.png",
            show_shapes=True,
            show_layer_names=True,
            )
```



## ▼ Test Model

```
ann.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
ann.fit(x_train, y_train, batch_size = 32, epochs = 100)
```

```

Epoch 1/100
366/366 [=====] - 3s 2ms/step - loss: 4548.2520 - accuracy: 0.0000e+00
Epoch 2/100
366/366 [=====] - 1s 2ms/step - loss: -3840.6453 - accuracy: 0.0000e+00
Epoch 3/100
366/366 [=====] - 1s 3ms/step - loss: -6026.0781 - accuracy: 0.0000e+00
Epoch 4/100
366/366 [=====] - 1s 4ms/step - loss: -7861.8940 - accuracy: 0.0000e+00
Epoch 5/100
366/366 [=====] - 1s 3ms/step - loss: -7861.9023 - accuracy: 0.0000e+00
Epoch 6/100
366/366 [=====] - 1s 3ms/step - loss: -7861.9092 - accuracy: 0.0000e+00
Epoch 7/100
366/366 [=====] - 1s 2ms/step - loss: -7861.9058 - accuracy: 0.0000e+00
Epoch 8/100
366/366 [=====] - 1s 2ms/step - loss: -7861.9102 - accuracy: 0.0000e+00
Epoch 9/100
366/366 [=====] - 1s 2ms/step - loss: -7861.9351 - accuracy: 0.0000e+00
Epoch 10/100
366/366 [=====] - 1s 2ms/step - loss: -7862.1113 - accuracy: 0.0000e+00
Epoch 11/100
366/366 [=====] - 1s 2ms/step - loss: -7862.1108 - accuracy: 0.0000e+00
Epoch 12/100
366/366 [=====] - 1s 2ms/step - loss: -7862.1216 - accuracy: 0.0000e+00
Epoch 13/100
366/366 [=====] - 1s 2ms/step - loss: -7862.1240 - accuracy: 0.0000e+00
Epoch 14/100
366/366 [=====] - 1s 2ms/step - loss: -7862.3057 - accuracy: 0.0000e+00
Epoch 15/100
366/366 [=====] - 1s 2ms/step - loss: -7862.3032 - accuracy: 0.0000e+00
Epoch 16/100

```

```
366/366 [=====] - 1s 2ms/step - loss: -7862.3066 - accuracy: 0.0000e+00
Epoch 17/100
366/366 [=====] - 1s 2ms/step - loss: -7862.3066 - accuracy: 0.0000e+00
Epoch 18/100
366/366 [=====] - 1s 2ms/step - loss: -7862.3066 - accuracy: 0.0000e+00
Epoch 19/100
366/366 [=====] - 1s 2ms/step - loss: -7862.3057 - accuracy: 0.0000e+00
Epoch 20/100
366/366 [=====] - 1s 3ms/step - loss: -7862.3066 - accuracy: 0.0000e+00
Epoch 21/100
366/366 [=====] - 1s 3ms/step - loss: -7862.3042 - accuracy: 0.0000e+00
Epoch 22/100
366/366 [=====] - 1s 3ms/step - loss: -7862.3057 - accuracy: 0.0000e+00
Epoch 23/100
366/366 [=====] - 1s 3ms/step - loss: -7862.3101 - accuracy: 0.0000e+00
Epoch 24/100
366/366 [=====] - 1s 2ms/step - loss: -7862.3042 - accuracy: 0.0000e+00
Epoch 25/100
366/366 [=====] - 1s 2ms/step - loss: -7862.3066 - accuracy: 0.0000e+00
Epoch 26/100
366/366 [=====] - 1s 2ms/step - loss: -7862.3042 - accuracy: 0.0000e+00
Epoch 27/100
366/366 [=====] - 1s 2ms/step - loss: -7862.3076 - accuracy: 0.0000e+00
Epoch 28/100
366/366 [=====] - 1s 2ms/step - loss: -7862.3091 - accuracy: 0.0000e+00
Epoch 29/100
366/366 [=====] - 1s 2ms/step - loss: -7862.3091 - accuracy: 0.0000e+00

y_pred = ann.predict(x_test)
y_pred = (y_pred > 0.5)
pd.DataFrame(list(zip(y_test, y_pred)), columns=['Actual', 'Predicted'])
```

92/92 [=====] - 1s 5ms/step

	Actual	Predicted
0	[1440]	[False, False, True, True, True, True, True, F...
1	[4270]	[False, False, True, True, True, True, True, F...
2	[1010]	[False, False, True, True, True, True, True, F...
3	[1970]	[False, False, True, True, True, True, True, F...
4	[2320]	[False, False, True, True, True, True, True, F...
...	...	...
2919	[4980]	[False, False, True, True, True, True, True, F...
2920	[2360]	[False, False, True, True, True, True, True, F...
2921	[2230]	[False, False, True, True, True, True, True, F...
2922	[2710]	[False, False, True, True, True, True, True, F...
2923	[1240]	[False, False, True, True, True, True, True, F...

2924 rows × 2 columns